Groq LPU Architecture and Specifications

The Groq Language Processing Unit (LPU) is a specialized AI hardware architecture.
It delivers deterministic execution without dynamic scheduling, reaching over 500 tokens/sec.
Model llama-3.1-8b-instant provides sub-second latency on Groq cloud infrastructure.

Page 2: Vector Search and ChromaDB

ChromaDB provides nearest neighbor indexing for embeddings generated from document chunk